

Contents

Chapter 1. Cover Page.....	1
Chapter 2. About this publication.....	2
IBM StoredIQ product library.....	2
Contacting IBM StoredIQ customer support.....	2
Chapter 3. IBM StoredIQ components.....	3
Solution components.....	3
Applications of IBM StoredIQ.....	4
Chapter 4. Planning for deployment.....	11
Open Virtual Appliance (OVA) configuration requirements.....	12

Chapter 1. Cover Page

Chapter 2. About this publication

IBM StoredIQ Deployment and Configuration Guide provides information about how to plan, deploy, and configure the IBM StoredIQ product.

IBM StoredIQ product library

The following documents are available in the IBM® StoredIQ® product library.

- *IBM StoredIQ Overview Guide*
- *IBM StoredIQ Deployment and Configuration Guide*
- *IBM StoredIQ Data Server Administration Guide*
- *IBM StoredIQ Administrator Administration Guide*
- *IBM StoredIQ Data Workbench User Guide*
- *IBM StoredIQ Policy Manager User Guide*
- *IBM StoredIQ Insights User Guide*
- *IBM StoredIQ Integration Guide*

Contacting IBM StoredIQ customer support

Chapter 3. IBM StoredIQ components

The IBM StoredIQ solution consists of these components: the application stack, the gateway, the data server, and optionally the Elasticsearch cluster.

Solution components

IBM StoredIQ provides three solution components: the gateway, data servers, and application stack (AppStack).

Gateway

The gateway communicates between the data servers and the application stack. The application stack polls the gateway for information about the data on the data servers. The data servers push the information to the gateway.

Data servers

A data server obtains the data from supported data sources and indexes it. By indexing this data, you gain information about unstructured data such as file size, file data types, file owners.

The data server pushes the information about volumes and indexes to the gateway so it can be communicated to the application stack. Multiple data servers feed into a single gateway.

Data servers can be categorized in two types: DataServer - Classic and DataServer - Distributed. A data server of the type DataServer - Classic uses the embedded PostgreSQL database for storing the index. With a data server of the type DataServer - Distributed, the index is stored in an Elasticsearch cluster. Data servers of this type also provide better performance in search queries. They can manage much larger amounts of data than data servers of the type DataServer - Classic, thus making the IBM StoredIQ deployments more scalable.

You can have both types of data servers in your IBM StoredIQ deployment.

In addition to completing standard administrative tasks, administrators can deploy the IBM StoredIQ Desktop Data Collector and index desktops from the data server.

Application stack

The application stack provides the user interface for the IBM StoredIQ Administrator, IBM StoredIQ Data Workbench, IBM StoredIQ Insights, and the IBM StoredIQ Policy Manager products.

The synchronization feature for integration with a governance catalog is also part of the application stack.

Elasticsearch cluster

The Elasticsearch cluster attached to a data server of the type DataServer - Distributed provides a single data store for all metadata and content of harvested objects. Indexed data is distributed automatically across the nodes in the cluster. Indexing and queries are load-balanced across all nodes. Nodes can be added dynamically without downtime and the indexing process can use these newly added nodes without further setup.

Applications of IBM StoredIQ

IBM StoredIQ provides interface applications that help fulfill its solution goals.

IBM StoredIQ Data Server

IBM StoredIQ Data Server user interface provides access to data server functionality. It allows administrators to view the dashboard and see the status of the jobs and system details. Administrators can manage information about servers and conduct various configurations on the system and application settings.

The screenshot shows the IBM StoredIQ Data Server user interface. The top navigation bar includes the IBM logo, a 'DS Admin' button, and tabs for 'Administration', 'Folders', and 'Audit'. Below this is a secondary bar with 'Dashboard', 'Data sources', and 'Configuration'. The main content area is divided into several sections:

- Page refresh:** Off | [30 sec](#) | [60 sec](#) | [90 sec](#)
- Today's job schedule:** No jobs scheduled for today.
- Jobs in progress:** View jobs as they run. No jobs are currently running.
- System summary:** View a summary of system details.

Total system data objects	10756
Total contained data objects	1081591
Total data objects	1092347
Number of volumes	13
Date of last completed harvest	No harvests run.
- Harvest statistics:** Review the performance over the last hour for all harvests.

Processes	4
Average data objects per second	0.0
Average data object size	0 bytes
Maximum data object size	0 bytes
Average data object processing time	0.0 sec
Maximum data object processing time	0.0 sec
- Event log:** The current event log as of 03/21/2018 05:03 PM. Includes links to 'Clear this view', 'Download today's event log', and 'View all event logs'. A list of 'Last 500 events' is shown, including database compactor completion, system maintenance, and audit cleaner actions.
- Appliance status:** Shows a green status indicator and a 'Controller' link. Below are links for 'About appliance' and 'View cache details'.

IBM StoredIQ Administrator

IBM StoredIQ Administrator helps you manage global assets common to the distributed infrastructure behind IBM StoredIQ applications.

The screenshot displays the IBM StoredIQ Administrator web interface. The top navigation bar includes the 'StoredIQ Administrator' title, a user profile 'super admin', a 'Help' link, and the IBM logo. A left-hand sidebar contains icons for various system components: Data, Data Servers, Volumes, Harvests, System Infosets, Users, Actions, Target Sets, Reports, Auto Classification, Cartridges, and Concepts.

The main content area is divided into several sections:

- All Data Currently Under Management:** A summary box showing 'Total Data Objects: 1,925,292', 'Total Data Size: 223.22 GB', 'Number of Data Servers: 2', and 'Number of Volumes: 196'.
- Search:** A search bar with the placeholder 'Enter key term(s)' and a 'Search' button.
- Buttons:** 'Add Volume' and 'View Volume(s)' buttons are located to the right of the search bar.
- Table:** A table with columns for 'Data server name', 'Status', 'IP Address', 'Data objects', and 'Total data object size'. It lists two data servers:

Data server name	Status	IP Address	Data objects	Total data object size
DS1	Healthy	192.168.224.114	1,135,057	135.11 GB
DS2	Healthy	192.168.225.179	790,235	88.11 GB
- Details: DS2:** A section for the selected data server, containing:
 - System Status:**
 - Status: Healthy
 - Status Message: System-and-services-running
 - IP Address: 192.168.225.179
 - Software Version: 7_6_0_14-STOREDIQ-8
 - Data Server Type: Classic
 - DB Version: 007.007.001.008
 - System Time: 16:53:44 +00:00
 - System Activity:**
 - Free RAM Memory: 10.43 GB of 15.58 GB
 - Free Swap Memory: 43.81 GB of 43.81 GB
 - Load Average: 0.00
 - Available Space: 2.01 TB
 - Active DB Connections: 0 of 512
 - System Uptime: 21 days, 6:01:52
- Buttons:** 'Restart Services' and 'Reboot Data Server' buttons are located to the right of the details section.

IBM StoredIQ Administrator provides at-a-glance understanding of the different issues that can crop up in the IBM StoredIQ environment. These views are unique to the IBM StoredIQ Administrator application as they provide an overview of how the system is running. They allow access to various pieces of information that are being shared across applications or allow for the management of resources in a centralized manner.

The administrator is the person responsible for managing the IBM StoredIQ. This individual has strong understanding of data sources, indexes, data servers, jobs, infosets, and actions. This list provides an overview as to how IBM StoredIQ Administrator works:

- Viewing data servers and volumes:** Using IBM StoredIQ Administrator, the Administrator can identify what data servers are deployed, their location, what data is being managed, and the status of each data server in the system. Volume management is a central component of IBM StoredIQ. IBM StoredIQ Administrator also allows the Administrator to see what volumes are currently under management, which data server is responsible for that volume, the state of the volume after indexing, and the amount and size of information that is contained by each volume. Administrators can also add volumes to and delete volumes from data servers through this interface.

If IBM StoredIQ is configured for integration with Information Governance Catalog, the Administrator can also manage which volumes are published to the governance catalog.

- **Scheduling harvests:** Harvesting, which can also be referred to as indexing, is the process or task by which IBM StoredIQ examines and classifies data in your network. Using IBM StoredIQ Administrator, harvests can be scheduled, edited, and deleted.
- **Creating system infosets:** System infosets that use only specific indexed volumes can be created and managed within IBM StoredIQ Administrator. Although infosets are a core component of IBM StoredIQ Data Workbench, system infosets are created as a shortcut for users in IBM StoredIQ Administrator.
- **Managing users:** The user management area allows administrators to create users and manage users' access to the various IBM StoredIQ applications.
- **Configuring and managing actions:** An action is any process that is taken upon the data that is represented by the indexes. Actions are run by data servers on indexed data objects. Any errors or warnings that are generated as a result of an action are recorded as exceptions in IBM StoredIQ Data Workbench.

Note: Actions can be created within IBM StoredIQ Administrator and then made available to other IBM StoredIQ applications such as IBM StoredIQ Data Workbench.

- **Managing target sets:** Provides an interface that allows the user to set the wanted targets for specific actions that require a destination volume for their actions.
- **Reports:** IBM StoredIQ Administrator provides a number of built-in reports, such as summaries of data objects in the system, storage use, and the number of identical documents in the system. You can create custom reports, including Query Analysis Reports for e-discovery purposes, and automatically email report notifications to administrators and other interested parties.
- **Auto-classification:** Automated document categorization, what IBM StoredIQ refers to as autotclassification models, integrates the IBM® Content Classification's classification model into the IBM StoredIQ infoset-generation process. Data Experts can use IBM Content Classification to train a classification model, which is then registered with IBM StoredIQ Administrator. The registered classification model can be applied to an existing infoset in IBM StoredIQ Data Workbench to generate new metadata for the objects in the infoset. Metadata can be used in rule-based filters to create new infosets.
- **Cartridges:** Cartridges are compressed files that contain analysis logic. When you add a cartridge to IBM StoredIQ AppStack, it can detect new data in documents during indexing and make these new insights searchable. For example, a sensitive pattern cartridge can enable IBM StoredIQ to detect passport numbers, phone numbers, and other IDs.

To apply the analysis logic contained in the cartridge, you must run a Step-up Analytics action that uses the cartridge on an infoset. IBM StoredIQ examines all documents in the infoset, applies the analytics, and then stores the analysis results in the IBM StoredIQ index.


- **Managing concepts:** Provides the ability to relate business concepts to indexed data.
- **Managing Mule scripts:** Helps you to create Mule scripts and upload script packages. These Mule scripts are used by IBM StoredIQ Policy Manager to create policies using the automation workflow.
- **DataSeter - Classic:** Data servers can be categorized in two types: DataSeter - Classic and DataSeter - Distributed. DataSeter - Classic refers to the regular data servers. It uses either the current PostgreSQL or Lucene index as an index.
- **DataSeter - Distributed:** The distributed data server uses an Elasticsearch cluster instead of an embedded Postgres database. It increases the scalability and flexibility of the IBM StoredIQ deployment in a way that it can manage much larger amounts of data. Without adding more data

servers, data that is managed by the IBM StoredIQ deployment can be increased by adding new nodes to the Elasticsearch cluster. Search queries perform better on DataServer - Distributed.

- **Connector API SDK:** A connector is a software component of IBM StoredIQ that is used to connect to a data source such as a network file system and access its data. Using IBM StoredIQ Connector API SDK, developers of other companies can develop connectors to new data sources outside the IBM StoredIQ development environment. These connectors can be integrated with a live IBM StoredIQ application to index, search, manage, and analyze data on the data source.

IBM StoredIQ Data Workbench

Big data is a pervasive problem, not a one-time occurrence. It is easy for most companies to realize that big data is problematic, but it is hard to identify what problems they have. Big data is all about the unknown, but the unknown cannot be off limits. IBM StoredIQ Data Workbench can help you learn about your data, make educated decisions with your most valuable asset, and turn your company's most dangerous risk into its most valuable asset.



Name	Total objects	Infoset size	Composition	Created	Type	Description
All Data Objects	1,925,292	223.22 GB	Mixed Level		System	All data objects.
All objects from SP (20108...	1,781	242.63 MB	Mixed Level	2015-12-13 11:44 AM	User	
All System-Level Objects	447,393	115.69 GB	Top Level		System	All system-level objects.
big12 ds2	423	37.92 MB	Mixed Level	2016-03-29 9:25 AM	System	
big12 ds2 user	423	37.92 MB	Mixed Level	2016-03-29 9:51 AM	User	
bmorgan-a ds1	4,273	5.22 GB	Mixed Level	2016-03-21 8:36 AM	System	
bmorgan-e ocr	57	160.11 MB	Top Level	2017-02-02 2:15 PM	System	
box2logesh	397	275.75 MB	Mixed Level	2016-03-22 11:40 AM	User	
bug 9168	17	157.96 MB	Top Level	2017-02-02 2:21 PM	User	
Collaborator Role Contains ...	46	14.52 MB	Top Level	2015-12-13 2:38 PM	User	
Collapsed - All objects from...	915	179.03 MB	Top Level	2015-12-13 11:47 AM	User	
DS1 > collaborator login na...	76	15.85 MB	Top Level	2015-12-13 2:28 PM	User	
DS1 all objects P8 nimmo8	58	3.28 MB	Mixed Level	2015-12-13 11:06 PM	User	

IBM StoredIQ Data Workbench is a data visualization and management tool that helps you to actively manage your company's data. It helps you to determine how much data you have, where it is, who owns it, and when it was last used. When you have a clear understanding of your company's data landscape, IBM StoredIQ Data Workbench helps you take control of data. You can make informed decisions about your data and act on that knowledge by copying, copying to retention, or conducting a discovery export.

Here are just some examples of how you can use IBM StoredIQ Data Workbench.

- You need to find all company email that is sent from or received by Eileen Sideways (esideways@thecompany.com). You can use IBM StoredIQ Data Workbench to find all email and then copy that data to a predefined repository. You can also use IBM StoredIQ Data Workbench to find all of the esideways@thecompany.com email that occurred between specific dates and then make that email available for review.
- As an administrator, you want to rid your networks and storage of unused data. You can use IBM StoredIQ Data Workbench to find all files that were not modified in more than five years.
- You want to find all image files that are created in 2007. Not only can IBM StoredIQ Data Workbench find all image files that were created in 2007. It also shows how much space they occupy on your network.
- A user needs to understand how data about Windows is being retained. Using IBM StoredIQ Data Workbench, you can provide that user with a visual overview of the number of objects that are retained and a breakdown of files per data source. Additionally, you can apply overlays to show the user if those files contain forbidden information such as credit-card numbers or Social Security numbers.
- If IBM StoredIQ is configured accordingly, you can select the info sets and filters that are published to the governance catalog for unified governance of structured and unstructured information. When integrating with Information Governance Catalog, you can also analyze and classify the data governed by IBM StoredIQ based on the data classes that are synchronized from the governance catalog.

IBM StoredIQ Insights

IBM StoredIQ Insights provides dynamic and interactive filtering for your data with easy access to all metadata and instant plain-text preview of document content for full-text indexed volumes.

Faceted search lets you drill down to refine your search results as needed. In addition, you can apply any valid IBM StoredIQ filter query. Tags let you categorize the data for easier management. Visual representations of search results help you gain further insights into your data. Several chart types let you look at and explore data from different perspectives, thus helping you identify patterns and relationships very quickly.

With IBM StoredIQ Insights, you can search data that is managed and indexed by a data server of the type DataServer - Distributed. In mixed deployments that have classic and distributed data servers, only the content from distributed data servers will be searchable.

IBM StoredIQ Insights

0_elastic User InfoSet: UNION(big12filetypes & big12filetypes-v2 & gdpv)

Facets

- Indexed annotation
- Tag
- Category
- File name extension
- Object size
- Created date
- Modified date
- Reviewer
- Error code
 - Unknown file format (Content e... (33)
 - Configured max. object size (Co... (1)
- Processing state
 - Content - Successful (23,724)
 - Content - Not performed (8,749)
 - Content - Failed (34)
 - Binary - Not performed (12,507)
 - System - Successful (12,507)
- Volume
- Path

Filter

Search

Save as InfoSet

Result list Basic charts Analytics charts 32,507 results

Update Tags

Object name	Size	Volume	Path	Created	Last modified
<input type="checkbox"/> rtfColoradoHealthCostInsuranceReport01172... Annual Report of the Commissioner of Insurance to The Colorado General Assembly on 2010 Hea...	1.25 MB	gdpv-es	MISC_documents/Industry/Insurance/Health I...	6/21/2017, 10:22:10 PM	8/3/2011, 1:41:08 AM Tags (5)
<input type="checkbox"/> CPT codes (HCPCS level 1).xls Sheet1 00164 ANESTH BIOPSY OF NOSE 00212ANESTH SKULL DRAINAGE ANESTH BLADDER TUMOR SURG	14.5 KB	big12filetypes-es	xistfiles	4/21/2014, 3:11:11 PM	1/1/2005, 5:36:34 PM Tags (5)
<input type="checkbox"/> Attachments, Binhex, Berlin.mail Berlin	27.55 KB	big12filetypes-v2-es	ost	1/10/2008, 6:01:00 PM	1/10/2008, 6:01:00 PM x-HUD, Insights, rev x
<input type="checkbox"/> Email with almost every attribute.mail	9.56 KB	big12filetypes-v2-es	ost	1/23/2008, 11:41:52 PM	1/23/2008, 11:41:52 ... x-exch x
<input type="checkbox"/> Certificate or license number.ppt Cert # 44939939	8 KB	big12filetypes-es	pptfiles	4/21/2014, 3:11:10 PM	1/1/2005, 7:36:34 PM Tags (9)
<input type="checkbox"/> Unicode text.mail chinese 汉语/漢語 hindi हिन्दी hebrew עברית german Abkürzungen Japanese 日本語 / 日本人の ...	7.73 KB	big12filetypes-v2-es	ost	1/31/2008, 11:16:36 PM	1/31/2008, 11:16:36 ...
<input type="checkbox"/> All attributes_amod.xls Sheet1 acct 333-222-333 acct number 434534534 Bill Jones was admitted on September 20, 1998 B...	22.5 KB	big12filetypes-es	xistfiles	4/21/2014, 3:11:11 PM	1/1/2005, 12:36:34 PM Tags (2)

IBM StoredIQ Policy Manager

IBM StoredIQ Policy Manager allows users to run mature policies and processes at scale across a wider range of data.

StoredIQ Policy Manager

super admin Help IBM

Policy Dashboard

View policy status. Click to view details.

Enter key terms(s) Search Create Policy

<input checked="" type="checkbox"/> Company A Created: 2018-03-21 11:07 AM On Off	Type: Automated Deletion When: 2018-03-21 11:12 AM What: DS1 Box My users Filter: Data Map: Fri Apr 01 2016 13:03:42 GMT-0500 (Central Sta...	Next Run: N/A Complete: 0 Admin: super admin Policy user: super admin
---	--	--

Loaded 1 of 1

The users can define and run systemwide policies, focusing on the execution of the process rather than understanding or reviewing affected data objects. Additionally, with reports of IBM StoredIQ Policy Manager, you can record what actions were conducted, when they were conducted, and what data was affected by the policy's execution.

IBM StoredIQ Desktop Data Collector

IBM StoredIQ Desktop Data Collector (also referred to as desktop client indexes desktops as volumes. The volumes appear in IBM StoredIQ Data Server and in IBM StoredIQ Administrator, where they can be used like any other data source.

The data server maintains an index using the information sent by the desktop client. After indexing, desktops - even offline or unreachable ones - can be viewed, searched, or targeted for later policy action.

Chapter 4. Planning for deployment

When you plan a deployment of IBM StoredIQ, evaluate several infrastructure considerations.

In addition to the information in this section, review the requirements detailed in the IBM Software Product Compatibility Reports (SPCR) tool at: [Software Product Compatibility Reports: StoredIQ 7.6](#)

If you plan to use IBM StoredIQ for Legal Identification and Collection to create and manage data boxes and data requests that are to be fulfilled by IBM StoredIQ, also check the system requirements for StoredIQ for Legal at: [Software Product Compatibility Reports: StoredIQ for Legal 2.0.3](#)

Generate customized reports with the SPCR tool

Go to the page at [Software Product Compatibility Reports](#) to create a high-level report for supported operating systems, related software, hypervisors, and supported translations for any product. You can also create an in-depth report to get detailed system requirements, hardware requirements, and end of service information for each product. You can search for a product in all of the report types and reports are generated based on your query values.

The following report types are the most commonly generated reports from software product compatibility reports:

Detailed system requirements

When you select your product version for the detailed system requirements report, you can set a report filter for **Operating system platforms**, **Product components**, and **Capabilities**, including prerequisites and support software. After you view the report, you can save it as a URL to generate anytime or download it as a PDF.

Hardware requirements

When you select your product version for the hardware requirements report, you can set a report filter by the **Operating system families** option. Set the operating system filter by selecting some or all of the operating systems that are supported by your product. After you view the report, you can save it as a URL to generate anytime or download it as a PDF.

End of service

The end of service report shows the service window of the products that you specify over an eightyyear span. For example, you can find out when your product is scheduled to go out of service.

Open Virtual Appliance (OVA) configuration requirements

IBM StoredIQ is deployed as virtual appliances and is supported in VMware ESXi 5.0 (all fix pack levels) or VMware ESXi 6.0 (all fix pack levels) environments. You must have a virtual infrastructure that meets the IBM StoredIQ hardware requirements.

Application stack

- vCPU: 2
- Memory: 4 GB
- Storage:
 - Primary disk (vmdisk1): 21 GB
 - Data disk (vmdisk2): 10 GB

Gateway server

- vCPU: 4
- Memory: 8 GB
- Storage:
 - Primary disk (vmdisk1): 100 GB
 - Data disk (vmdisk2): 75 GB
 - Swap disk (vmdisk3): 40 - 100 GB

DataServer - Classic


- vCPU: 4

Even though increasing the number of vCPUs increases performance, the actual benefits depend on whether the specific host is oversubscribed or not.
- Memory: 16 GB

While the minimum value works under light-load condition, as the load increases, the data server quickly starts using swap space. For high load situations, increasing RAM beyond 16 GB can benefit performance. Monitoring swap usage can provide insight.
- Storage:
 - Primary disk (vmdisk1, SCSI 0:0): Default is 150 GB

This virtual disk has an associated VMDK that contains the IBM StoredIQ operating code. Do not change its size.

Attention

 **Attention:** If you delete the primary disk, you delete the operating system, and the IBM StoredIQ software; the virtual machine might need to be scrapped.

– Data disk (vmdisk2, SCSI 0:1): Default is 2 TB

This virtual disk can be resized according to expectations on the amount of harvest data to be stored. For purposes of estimation, the index storage requirement for metadata is about 30 GB per TB of managed source data. Full-text indexing requires an extra 170 GB per TB. Therefore, the default data disk size is targeted for managing 10 TB of source information.

– Swap disk (vmdisk3, SCSI 0:2): Default is 40 GB

When under load, the data server can use much RAM; therefore, having ample swap space is prudent. The minimum swap size is equal to the amount of RAM configured for the virtual machine. For best performance under load, place this disk on the highest speed data store available to the host.

The general size limits for a data server are 150 million objects or 500 defined volumes, whichever limit is reached first. Assuming an average object size of 200 KB equals about 30 TB of managed storage across 30 volumes of 5 million objects each, the index storage requirement for metadata on ~30 TB of storage that contains uncompressed general office documents is ~330 GB (11 GB per TB). Add 100 GB per TB of managed storage for full-text or snippet index. For example, to support 30 TB of storage that is indexed for metadata, you need 8 TB indexed for full-text search and extracted text (snippet cache) of 8 TB for auto-classification. A total of 1.9 TB of storage is required (metadata: 330 GB, full-text: 800 GB, snippet cache: 800 GB).

Data-server performance is impacted by the IOPS available from the storage subsystem. For each data server under maximum workload, at least 650 IOPS generally delivers acceptable performance. In the situations where there is a high load on the system, the IOPS that is used can reach up to 7000 with main write operations.

DataSeter - Distributed

- vCPU: 4
- Memory: 16 GB
- Storage:
 - Primary disk (vmdisk1, SCSI 0:0): Default is 150 GB
 - Data disk (vmdisk2, SCSI 0:1): Default is 2 TB
 - Swap disk (vmdisk3, SCSI 0:2): Default is 40 GB

If you deploy this type of data server, you must also deploy an Elasticsearch cluster with at least one node. If you deploy a cluster with more nodes, each of the Elasticsearch nodes must meet the listed requirements.

Each Elasticsearch node

- vCPU: 1
- Memory: 32 GB
- Storage:
 - Primary disk (vmdisk1): 100 GB
 - Data disk (vmdisk2): 1 TB

The required memory depends on the index size that is handled by the data node. For a better performance, consider increasing the memory. The total memory available to the node must be distributed between the host operating system, the Elasticsearch java heap space, and the kernel file system caches. For example, if the system has 32 GB memory, 2 GB must be allocated for the host operating system, 15 GB for the java heap space, and the remaining 15 GB for the file system caches.

The data disk (vmdisk2) can be resized according to expectations on the amount of harvest data to be stored.